

# Artificial Intelligence Nanodegree

## Convolutional Neural Networks

### Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

---

### Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an

estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).



In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

## The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- [Step 0](#): Import Datasets
- [Step 1](#): Detect Humans
- [Step 2](#): Detect Dogs
- [Step 3](#): Create a CNN to Classify Dog Breeds (from Scratch)
- [Step 4](#): Use a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 5](#): Create a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 6](#): Write your Algorithm
- [Step 7](#): Test Your Algorithm

## Step 0: Import Datasets

### Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded

classification labels

- dog\_names - list of string-valued dog breed names for translating labels

In [1]:

```
1  from sklearn.datasets import load_files
2  from keras.utils import np_utils
3  import numpy as np
4  from glob import glob
5
6  # define function to load train, test, and validation datasets
7  def load_dataset(path):
8      data = load_files(path)
9      dog_files = np.array(data['filenames'])
10     dog_targets = np_utils.to_categorical(np.array(data['target']), 133)
11     return dog_files, dog_targets
12
13 # load train, test, and validation datasets
14 train_files, train_targets = load_dataset('dogImages/train')
15 valid_files, valid_targets = load_dataset('dogImages/valid')
16 test_files, test_targets = load_dataset('dogImages/test')
17
18 # load list of dog names
19 dog_names = [item[20:-1] for item in sorted(glob("dogImages/train/*/"))]
20
21 # print statistics about the dataset
22 print('There are %d total dog categories.' % len(dog_names))
23 print('There are %s total dog images.\n' % len(np.hstack([train_files, valid_files, test_files])))
24 print('There are %d training dog images.' % len(train_files))
25 print('There are %d validation dog images.' % len(valid_files))
26 print('There are %d test dog images.' % len(test_files))
```

```
/Users/Anna/anaconda/envs/py27/lib/python2.7/site-packages/h5py/__init__
__.py:36: FutureWarning: Conversion of the second argument of issubdtype
pe from `float` to `np.floating` is deprecated. In future, it will be
treated as `np.float64 == np.dtype(float).type`.
```

```
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

There are 133 total dog categories.

There are 8351 total dog images.

There are 6680 training dog images.

There are 835 validation dog images.

There are 836 test dog images.

## Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array `human_files`.

In [2]:

```
1 import random
2 random.seed(8675309)
3
4 # load filenames in shuffled human dataset
5 human_files = np.array(glob("lfw/**/*.jpg"))
6 random.shuffle(human_files)
7
8 # print statistics about the dataset
9 print('There are %d total human images.' % len(human_files))
```

There are 13233 total human images.

## Step 1: Detect Humans


We use OpenCV's implementation of [Haar feature-based cascade classifiers](http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html) ([http://docs.opencv.org/trunk/d7/d8b/tutorial\\_py\\_face\\_detection.html](http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html)) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on [github](https://github.com/opencv/opencv/tree/master/data/haarcascades) (<https://github.com/opencv/opencv/tree/master/data/haarcascades>). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

In [3]:

```
1 import cv2
2 import matplotlib.pyplot as plt
3 %matplotlib inline
4
5 # extract pre-trained face detector
6 face_cascade = cv2.CascadeClassifier('haarcascades/haarcascade_frontalface_alt.xml')
7
8 # load color (BGR) image
9 img = cv2.imread(human_files[3])
10 # convert BGR image to grayscale
11 gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
12
13 # find faces in image
14 faces = face_cascade.detectMultiScale(gray)
15
16 # print number of faces detected in the image
17 print('Number of faces detected:', len(faces))
18
19 # get bounding box for each detected face
20 for (x,y,w,h) in faces:
21     # add bounding box to color image
22     cv2.rectangle(img,(x,y),(x+w,y+h),(255,0,0),2)
23
24 # convert BGR image to RGB for plotting
25 cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
26
27 # display the image, along with bounding box
28 plt.imshow(cv_rgb)
29 plt.show()
```

('Number of faces detected:', 1)



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

## Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

In [4]:

```
1  # returns "True" if face is detected in image stored at img_path
2  def face_detector(img_path):
3      img = cv2.imread(img_path)
4      gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
5      faces = face_cascade.detectMultiScale(gray)
6      return len(faces) > 0
```

## (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:**

Humans face detection accuracy: 99%

Dogs face detection accuracy: 11%

In [5]:

```
1 human_files_short = human_files[:100]
2 dog_files_short = train_files[:100]
3 # Do NOT modify the code above this line.
4
5 ## TODO: Test the performance of the face_detector algorithm
6 ## on the images in human_files_short and dog_files_short.
7 human_pred = []
8 for file in human_files_short: human_pred.append(int(face_detector(file)))
9 human_score = np.mean(human_pred)
10 print ('Humans face detection accuracy:', human_score)
11
12 dog_pred = []
13 for file in dog_files_short: dog_pred.append(int(face_detector(file)))
14 dog_score = np.mean(dog_pred)
15 print('Dogs face detection accuracy:', dog_score)
16
```

```
('Humans face detection accuracy:', 0.99)
('Dogs face detection accuracy:', 0.11)
```

**Question 2:** This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unnecessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

**Answer:**

Yes, I think it is a reasonable expectation to pose to the user that we accept human images only when they provide a clear view of a face. In image detection task, there are many features OpenCV could be looking for to classify it as human. For example the boundaries of a human face, it would be hard to correctly classify it if the image of the human face is blurry or covered. The higher the quality of the image helps the algorithm to be more efficient with higher accuracy.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.

In [6]:

```
1 ## (Optional) TODO: Report the performance of another
2 ## face detection algorithm on the LFW dataset
3 ### Feel free to use as many code cells as needed.
```



## Step 2: Detect Dogs

In this section, we use a pre-trained [ResNet-50](http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006) (<http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006>) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on [ImageNet](http://www.image-net.org/) (<http://www.image-net.org/>), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

In [7]:

```
1 from keras.applications.resnet50 import ResNet50
2
3 # define ResNet50 model
4 ResNet50_model = ResNet50(weights='imagenet')
```



## Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

$$(nb\_samples, rows, columns, channels),$$
$$(nb\_samples, rows, columns, channels),$$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is  $224 \times 224$  pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

$$(1, 224, 224, 3).$$
$$(1, 224, 224, 3).$$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

$$(nb\_samples, 224, 224, 3).$$
$$(nb\_samples, 224, 224, 3).$$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!

In [8]:

```
1  from keras.preprocessing import image
2  from tqdm import tqdm
3
4  def path_to_tensor(img_path):
5      # loads RGB image as PIL.Image.Image type
6      img = image.load_img(img_path, target_size=(224, 224))
7      # convert PIL.Image.Image type to 3D tensor with shape (224, 224, 3)
8      x = image.img_to_array(img)
9      # convert 3D tensor to 4D tensor with shape (1, 224, 224, 3) and return 4D
10     return np.expand_dims(x, axis=0)
11
12 def paths_to_tensor(img_paths):
13     list_of_tensors = [path_to_tensor(img_path) for img_path in tqdm(img_paths)]
14     return np.vstack(list_of_tensors)
```

## Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as `[103.939, 116.779, 123.68]`) and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` [here](https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py) ([https://github.com/fchollet/keras/blob/master/keras/applications/imagenet\\_utils.py](https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py)).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose  $ii$ -th entry is the model's predicted probability that the image belongs to the  $ii$ -th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the `argmax` of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this [dictionary](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>).

In [9]:

```
1 from keras.applications.resnet50 import preprocess_input, decode_predictions
2
3 def ResNet50_predict_labels(img_path):
4     # returns prediction vector for image located at img_path
5     img = preprocess_input(path_to_tensor(img_path))
6     return np.argmax(ResNet50_model.predict(img))
```

## Write a Dog Detector

While looking at the [dictionary](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the `ResNet50_predict_labels` function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

In [10]:

```
1 """ returns "True" if a dog is detected in the image stored at img_path
2 def dog_detector(img_path):
3     prediction = ResNet50_predict_labels(img_path)
4     return ((prediction <= 268) & (prediction >= 151))
```

## (IMPLEMENTATION) Assess the Dog Detector

**Question 3:** Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

**Answer:**

Humans face detection accuracy: 2%

Dogs face detection accuracy: 100%

In [11]:

```
1  ### TODO: Test the performance of the dog_detector function
2  ### on the images in human_files_short and dog_files_short.
3
4  human_pred = []
5  for file in human_files_short: human_pred.append(int(dog_detector(file)))
6  human_score = np.mean(human_pred)
7  print ('Humans face detection accuracy:', human_score)
8
9  dog_pred = []
10 for file in dog_files_short: dog_pred.append(int(dog_detector(file)))
11 dog_score = np.mean(dog_pred)
12 print ('Dogs face detection accuracy:', dog_score)
13
14
```

```
('Humans face detection accuracy:', 0.02)
('Dogs face detection accuracy:', 1.0)
```


## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning yet!), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

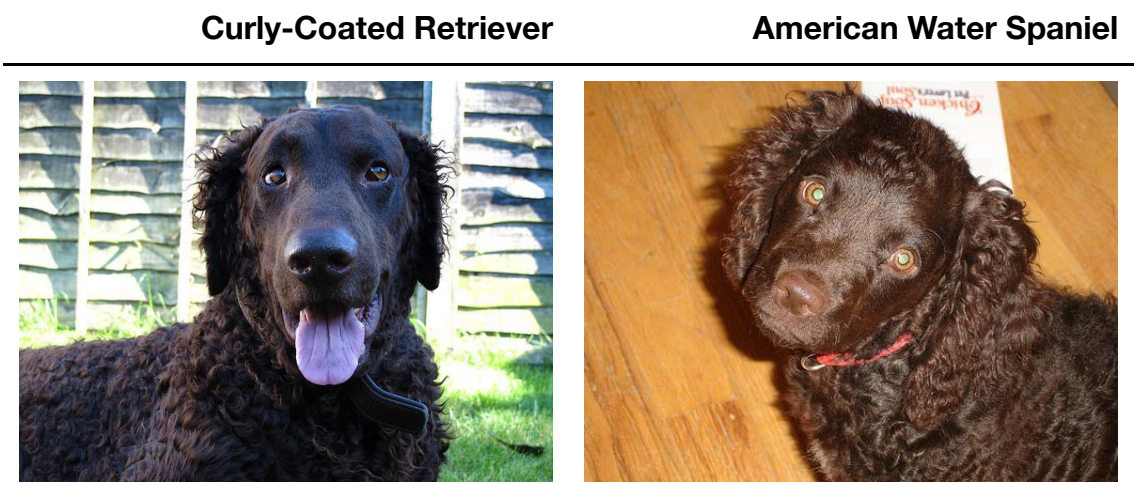
Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this estimate to figure out how long it will take for your algorithm to train.



We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.



It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).



Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.



We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imabalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

## Pre-process the Data

We rescale the images by dividing every pixel in every image by 255.

In [12]:

```
1  from PIL import ImageFile
2  ImageFile.LOAD_TRUNCATED_IMAGES = True
3
4  # pre-process the data for Keras
5  train_tensors = paths_to_tensor(train_files).astype('float32')/255
6  valid_tensors = paths_to_tensor(valid_files).astype('float32')/255
7  test_tensors = paths_to_tensor(test_files).astype('float32')/255
```

```
100%|██████████| 6680/6680 [02:33<00:00, 43.39it/s]
100%|██████████| 835/835 [00:17<00:00, 48.52it/s]
100%|██████████| 836/836 [00:17<00:00, 48.48it/s]
```

## (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:

| Layer (type)  | Output Shape         | Param # |
|---|----------------------|---------|
| conv2d_1 (Conv2D)                                   | (None, 223, 223, 16) | 208     |
| max_pooling2d_1 (MaxPooling2D)                      | (None, 111, 111, 16) | 0       |
| conv2d_2 (Conv2D)                                   | (None, 110, 110, 32) | 2080    |
| max_pooling2d_2 (MaxPooling2D)                      | (None, 55, 55, 32)   | 0       |
| conv2d_3 (Conv2D)                                   | (None, 54, 54, 64)   | 8256    |
| max_pooling2d_3 (MaxPooling2D)                      | (None, 27, 27, 64)   | 0       |
| global_average_pooling2d_1 (GlobalAveragePooling2D) | (None, 64)           | 0       |
| dense_1 (Dense)                                     | (None, 133)          | 8645    |
| Total params: 19,189.0                              |                      |         |
| Trainable params: 19,189.0                          |                      |         |
| Non-trainable params: 0.0                           |                      |         |



**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

**Answer:**

I chose a Convolutional layer of 16 filters to handle the initial input of the Neural Network. I also chose the activation function relu. I followed this layer with a max pooling layer of 2 which is standard.

Added another convolutional layer that is double the size as the first layer to get more detailed features as passed through from the prior layer. Nothing was changed here as it was designed purposely to be the exact same.

Following the first 2 layers I changed the third convolutional layer to include 64 filters for more details. This way I could give the network more information to classify the dog breed successfully. With the rest remain the same.

To avoid overfitting, I added dropout layer followed by a flatten layer so I could then add some fully connected layers to the network and dropout layer again. Then rape up with a output layer with softmax function.

In [13]:

```

1  from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D
2  from keras.layers import Dropout, Flatten, Dense
3  from keras.models import Sequential
4
5  model = Sequential()
6
7  ### TODO: Define your architecture.
```



```
8 model.add(Conv2D(filters=16, kernel_size=2, padding='same', activation='relu',
9                  input_shape=(224, 224, 3)))
10 model.add(MaxPooling2D(pool_size=2))
11 model.add(Conv2D(filters=32, kernel_size=2, padding='same', activation='relu'))
12 model.add(MaxPooling2D(pool_size=2))
13 model.add(Conv2D(filters=64, kernel_size=2, padding='same', activation='relu'))
14 model.add(MaxPooling2D(pool_size=2))
15
16 model.add(Dropout(0.3))
17 model.add(Flatten())
18 model.add(Dense(500, activation='relu'))
19 model.add(Dropout(0.4))
20 model.add(Dense(133, activation='softmax'))
21
22
23 model.summary()
```

| Layer (type)                   | Output Shape         | Param #  |
|--------------------------------|----------------------|----------|
| =====                          |                      |          |
| conv2d_1 (Conv2D)              | (None, 224, 224, 16) | 208      |
| <hr/>                          |                      |          |
| max_pooling2d_2 (MaxPooling2D) | (None, 112, 112, 16) | 0        |
| <hr/>                          |                      |          |
| conv2d_2 (Conv2D)              | (None, 112, 112, 32) | 2080     |
| <hr/>                          |                      |          |
| max_pooling2d_3 (MaxPooling2D) | (None, 56, 56, 32)   | 0        |
| <hr/>                          |                      |          |
| conv2d_3 (Conv2D)              | (None, 56, 56, 64)   | 8256     |
| <hr/>                          |                      |          |
| max_pooling2d_4 (MaxPooling2D) | (None, 28, 28, 64)   | 0        |
| <hr/>                          |                      |          |
| dropout_1 (Dropout)            | (None, 28, 28, 64)   | 0        |
| <hr/>                          |                      |          |
| flatten_2 (Flatten)            | (None, 50176)        | 0        |
| <hr/>                          |                      |          |
| dense_1 (Dense)                | (None, 500)          | 25088500 |
| <hr/>                          |                      |          |
| dropout_2 (Dropout)            | (None, 500)          | 0        |
| <hr/>                          |                      |          |
| dense_2 (Dense)                | (None, 133)          | 66633    |
| =====                          |                      |          |
| Total params: 25,165,677       |                      |          |
| Trainable params: 25,165,677   |                      |          |
| Non-trainable params: 0        |                      |          |
| <hr/>                          |                      |          |

## Compile the Model



In [14]:

```
1 model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc
```

## (IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

In [15]:

```
1 from keras.callbacks import ModelCheckpoint
2
3 ### TODO: specify the number of epochs that you would like to use to train the model
4
5 epochs = 5
6
7 ### Do NOT modify the code below this line.
8
9 checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.from_scratch.hdf5',
10                                verbose=1, save_best_only=True)
11
12 model.fit(train_tensors, train_targets,
13           validation_data=(valid_tensors, valid_targets),
14           epochs=epochs, batch_size=20, callbacks=[checker], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/5

6660/6680 [=====>.] - ETA: 1s - loss: 5.0752 - acc: 0.0207

Epoch 00001: val\_loss improved from inf to 4.58885, saving model to saved\_models/weights.best.from\_scratch.hdf5

6680/6680 [=====] - 428s 64ms/step - loss: 5.0724 - acc: 0.0208 - val\_loss: 4.5888 - val\_acc: 0.0359

Epoch 2/5

6660/6680 [=====>.] - ETA: 1s - loss: 4.3535 - acc: 0.0590

Epoch 00002: val\_loss improved from 4.58885 to 4.24946, saving model to saved\_models/weights.best.from\_scratch.hdf5

6680/6680 [=====] - 466s 70ms/step - loss: 4.3521 - acc: 0.0594 - val\_loss: 4.2495 - val\_acc: 0.0647

Epoch 3/5

6660/6680 [=====>.] - ETA: 1s - loss: 3.7908 - acc: 0.1380

Epoch 00003: val\_loss improved from 4.24946 to 4.09295, saving model to saved\_models/weights.best.from\_scratch.hdf5

6680/6680 [=====] - 464s 70ms/step - loss: 3.7916 - acc: 0.1379 - val\_loss: 4.0929 - val\_acc: 0.0826

```
Epoch 4/5
6660/6680 [=====>.] - ETA: 1s - loss: 2.9790 -
acc: 0.2886
Epoch 00004: val_loss did not improve
6680/6680 [=====] - 472s 71ms/step - loss: 2.
9797 - acc: 0.2886 - val_loss: 4.2988 - val_acc: 0.0898
Epoch 5/5
6660/6680 [=====>.] - ETA: 1s - loss: 1.9469 -
acc: 0.5159
Epoch 00005: val_loss did not improve
6680/6680 [=====] - 411s 62ms/step - loss: 1.
9477 - acc: 0.5159 - val_loss: 4.9279 - val_acc: 0.0898
```

Out[15]:

<keras.callbacks.History at 0x11c23a450>

## Load the Model with the Best Validation Loss

In [16]:

```
1 model.load_weights('saved_models/weights.best.from_scratch.hdf5')
```

## Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

In [17]:

```
1 # get index of predicted dog breed for each image in test set
2 dog_breed_predictions = [np.argmax(model.predict(np.expand_dims(tensor, axis=0))
3
4 # report test accuracy
5 test_accuracy = 100*np.sum(np.array(dog_breed_predictions)==np.argmax(test_target
6 print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 8.0000%

## Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

## Obtain Bottleneck Features

In [18]:

```
1 bottleneck_features = np.load('bottleneck_features/DogVGG16Data.npz')
2 train_VGG16 = bottleneck_features['train']
3 valid_VGG16 = bottleneck_features['valid']
4 test_VGG16 = bottleneck_features['test']
```

## Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

In [19]:

```
1 VGG16_model = Sequential()
2 VGG16_model.add(GlobalAveragePooling2D(input_shape=train_VGG16.shape[1:]))
3 VGG16_model.add(Dense(133, activation='softmax'))
4
5 VGG16_model.summary()
```

| Layer (type)                             | Output Shape | Param # |
|--|--------------|---------|
| =====                                    |              |         |
| global_average_pooling2d_1 ( (None, 512) |              | 0       |
| dense_3 (Dense)                          | (None, 133)  | 68229   |
| =====                                    |              |         |
| Total params: 68,229                     |              |         |
| Trainable params: 68,229                 |              |         |
| Non-trainable params: 0                  |              |         |

## Compile the Model

In [20]:

```
1 VGG16_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
```

## Train the Model

In [21]:

```
1  checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.VGG16.hdf5',
2                                  verbose=1, save_best_only=True)
3
4  VGG16_model.fit(train_VGG16, train_targets,
5                  validation_data=(valid_VGG16, valid_targets),
6                  epochs=20, batch_size=20, callbacks=[checker], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/20

6580/6680 [=====>.] - ETA: 0s - loss: 12.3766 - acc: 0.1309

Epoch 00001: val\_loss improved from inf to 10.68090, saving model to saved\_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 4s 633us/step - loss: 12.3627 - acc: 0.1316 - val\_loss: 10.6809 - val\_acc: 0.2287

Epoch 2/20

6540/6680 [=====>.] - ETA: 0s - loss: 10.1996 - acc: 0.2872

Epoch 00002: val\_loss improved from 10.68090 to 9.86518, saving model to saved\_models/weights.best.VGG16.hdf5

6680/6680 [=====] - 2s 261us/step - loss: 10.1736 - acc: 0.2888 - val\_loss: 9.8652 - val\_acc: 0.3150

Epoch 3/20

6580/6680 [=====>.] - ETA: 0s - loss: 9.5918 - acc: 0.3457

Epoch 00003: val\_loss improved from 9.86518 to 9.58232, saving model to saved\_models/weights.best.VGG16.hdf5

## Load the Model with the Best Validation Loss

In [22]:

```
1  VGG16_model.load_weights('saved_models/weights.best.VGG16.hdf5')
```

## Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

In [23]:

```
1  # get index of predicted dog breed for each image in test set
2  VGG16_predictions = [np.argmax(VGG16_model.predict(np.expand_dims(feature, axis=
3
4  # report test accuracy
5  test_accuracy = 100*np.sum(np.array(VGG16_predictions)==np.argmax(test_targets,
6  print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 47.0000%

## Predict Dog Breed with the Model

In [24]:

```
1  from extract_bottleneck_features import *
2
3  def VGG16_predict_breed(img_path):
4      # extract bottleneck features
5      bottleneck_feature = extract_VGG16(path_to_tensor(img_path))
6      # obtain predicted vector
7      predicted_vector = VGG16_model.predict(bottleneck_feature)
8      # return dog breed that is predicted by the model
9      return dog_names[np.argmax(predicted_vector)]
```

## Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- [VGG-19 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz) bottleneck features
- [ResNet-50 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz) bottleneck features
- [Inception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz) bottleneck features
- [Xception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where `{network}`, in the above filename, can be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

### (IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```
bottleneck_features = np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']
```

In [25]:

```
1  ### TODO: Obtain bottleneck features from another pre-trained CNN.
2  bottleneck_features = np.load('bottleneck_features/DogResNet50Data.npz')
3  train_DogResNet50 = bottleneck_features['train']
4  valid_DogResNet50 = bottleneck_features['valid']
5  test_DogResNet50 = bottleneck_features['test']
```

## (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:**

I decided to leverage the current resnet weights and make that an input into the global average pooling layer to take advantage of transfer learning. So I could make use of what the network already understood from previous training. I added the fully connected layer with an output of 133 units, because there are 133 dognames, with the softmax function.

In [26]:

```
1  ### TODO: Define your architecture.
2  Resnet50_model = Sequential()
3  Resnet50_model.add(GlobalAveragePooling2D(input_shape=train_DogResNet50.shape[1:]))
4  Resnet50_model.add(Dense(133, activation='softmax'))
5
6  Resnet50_model.summary()
7
```

| Layer (type)  | Output Shape | Param # |
|---|--------------|---------|
| =====   |              |         |
| global_average_pooling2d_2 (GlobalAveragePooling2D) | (None, 2048) | 0       |
| dense_4 (Dense)                                     | (None, 133)  | 272517  |
| =====   |              |         |
| Total params: 272,517                               |              |         |
| Trainable params: 272,517                           |              |         |
| Non-trainable params: 0                             |              |         |

## (IMPLEMENTATION) Compile the Model

In [27]:

```
1  ### TODO: Compile the model.
2  Resnet50_model.compile(loss='categorical_crossentropy', optimizer='rmsprop', metrics=['accuracy'])
3
```



## (IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

In [28]:

```
1  ### TODO: Train the model.
2
3  checkpointer = ModelCheckpoint(filepath='saved_models/weights.best.ResNet50.hdf5',
4                                  verbose=1, save_best_only=True)
5
6  Resnet50_model.fit(train_DogResNet50, train_targets,
7                    validation_data=(valid_DogResNet50, valid_targets),
8                    epochs=10, batch_size=20, callbacks=[checker], verbose=1)
```

Train on 6680 samples, validate on 835 samples

Epoch 1/10

6560/6680 [=====>.] - ETA: 0s - loss: 1.6179 - acc: 0.6003

Epoch 00001: val\_loss improved from inf to 0.83293, saving model to saved\_models/weights.best.ResNet50.hdf5

6680/6680 [=====] - 3s 499us/step - loss: 1.6011 - acc: 0.6037 - val\_loss: 0.8329 - val\_acc: 0.7425

Epoch 2/10

6660/6680 [=====>.] - ETA: 0s - loss: 0.4435 - acc: 0.8614

Epoch 00002: val\_loss improved from 0.83293 to 0.71584, saving model to saved\_models/weights.best.ResNet50.hdf5

6680/6680 [=====] - 2s 348us/step - loss: 0.4431 - acc: 0.8614 - val\_loss: 0.7158 - val\_acc: 0.7904

Epoch 3/10

6660/6680 [=====>.] - ETA: 0s - loss: 0.2663 - acc: 0.9179

Epoch 00003: val\_loss improved from 0.71584 to 0.66617, saving model to saved\_models/weights.best.ResNet50.hdf5

6680/6680 [=====] - 2s 347us/step - loss: 0.2663 - acc: 0.9178 - val\_loss: 0.6662 - val\_acc: 0.8072

Epoch 4/10

6580/6680 [=====>.] - ETA: 0s - loss: 0.1816 - acc: 0.9418

Epoch 00004: val\_loss improved from 0.66617 to 0.66370, saving model to saved\_models/weights.best.ResNet50.hdf5

6680/6680 [=====] - 2s 338us/step - loss: 0.1808 - acc: 0.9416 - val\_loss: 0.6637 - val\_acc: 0.8072

Epoch 5/10

6600/6680 [=====>.] - ETA: 0s - loss: 0.1236 - acc: 0.9636

Epoch 00005: val\_loss did not improve

```
6680/6680 [=====] - 2s 342us/step - loss: 0.1
230 - acc: 0.9638 - val_loss: 0.7005 - val_acc: 0.8000
Epoch 6/10
6640/6680 [=====>.] - ETA: 0s - loss: 0.0926 -
acc: 0.9723
Epoch 00006: val_loss did not improve
6680/6680 [=====] - 2s 368us/step - loss: 0.0
922 - acc: 0.9725 - val_loss: 0.7341 - val_acc: 0.8084
Epoch 7/10
6660/6680 [=====>.] - ETA: 0s - loss: 0.0691 -

acc: 0.9802
Epoch 00007: val_loss did not improve
6680/6680 [=====] - 3s 378us/step - loss: 0.0
690 - acc: 0.9802 - val_loss: 0.7082 - val_acc: 0.8180
Epoch 8/10
6560/6680 [=====>.] - ETA: 0s - loss: 0.0496 -
acc: 0.9858
Epoch 00008: val_loss did not improve
6680/6680 [=====] - 2s 368us/step - loss: 0.0
500 - acc: 0.9856 - val_loss: 0.6813 - val_acc: 0.8335
Epoch 9/10
6580/6680 [=====>.] - ETA: 0s - loss: 0.0371 -
acc: 0.9895
Epoch 00009: val_loss did not improve
6680/6680 [=====] - 2s 334us/step - loss: 0.0
371 - acc: 0.9895 - val_loss: 0.7490 - val_acc: 0.8204
Epoch 10/10
6520/6680 [=====>.] - ETA: 0s - loss: 0.0279 -
acc: 0.9926
Epoch 00010: val_loss did not improve
6680/6680 [=====] - 2s 331us/step - loss: 0.0
276 - acc: 0.9928 - val_loss: 0.7553 - val_acc: 0.8204
```

Out[28]:

<keras.callbacks.History at 0x11d259d50>

## (IMPLEMENTATION) Load the Model with the Best Validation Loss

In [29]:

```
1  ### TODO: Load the model weights with the best validation loss.
2  Resnet50_model.load_weights('saved_models/weights.best.ResNet50.hdf5')
3
```

## (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

In [30]:

```
1  ### TODO: Calculate classification accuracy on the test dataset.
2  # get index of predicted dog breed for each image in test set
3  Resnet50_predictions = [np.argmax(Resnet50_model.predict(np.expand_dims(feature
4
5  # report test accuracy
6  test_accuracy = 100*np.sum(np.array(Resnet50_predictions)==np.argmax(test_target
7  print('Test accuracy: %.4f%%' % test_accuracy)
```

Test accuracy: 82.0000%

## (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed ( Affenpinscher , Afghan\_hound , etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the argmax of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py` , and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the function

`extract_{network}`

where `{network}` , in the above filename, should be one of `VGG19` , `Resnet50` , `InceptionV3` , or `Xception` .

In [31]:

```
1  ### TODO: Write a function that takes a path to an image as input
2  ### and returns the dog breed that is predicted by the model.
3  def ResNet50_predict_breed(img_path):
4      # extract bottleneck features
5      bottleneck_feature = extract_Resnet50(path_to_tensor(img_path))
6      # obtain predicted vector
7      predicted_vector = Resnet50_model.predict(bottleneck_feature)
8      # return dog breed that is predicted by the model
9      return dog_names[np.argmax(predicted_vector)]
10
```

## Step 6: Write your Algorithm


Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

```
hello, human!
```



```
You look like a ...  
Chinese_shar-pei
```

## (IMPLEMENTATION) Write your Algorithm

In [32]:

```
1  ### TODO: Write your algorithm.
2  ### Feel free to use as many code cells as needed.
3
4  def dog_breed_detector(img_path):
5      breed = ResNet50_predict_breed(img_path)
6
7      # Display the image
8      img = cv2.imread(img_path)
9      cv_rgb = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
10     plt.imshow(cv_rgb)
11     plt.show()
12
13     # Detect what it is
14     if dog_detector(img_path):
15         print("That's a dog in the image. Breed: " + str(breed))
16     elif face_detector(img_path):
17         print("That's a human in the image, but it looks like dog breed " + str
18     else:
19         print("I don't know what's in the image.")
20
21     dog_breed_detector(train_files[0])
```

A local file was found, but it seems to be incomplete or outdated because the md5 file hash does not match the original value of a268eb855778b3df3c7506639542a6af so we will re-download the data.

Downloading data from [https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5) (https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5)

94658560/94653016 [=====] - 53s 1us/step

94666752/94653016 [=====] - 53s 1us/step



That's a dog in the image. Breed: Kuvasz

## Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

### (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

#### Answer:

The output is better than I expected since all 6 images are correctly classified between human and dog, also dog breed are all correctly classified! Even in picture2 Bullmastiff were wearing clothes, the detector still correctly classified it. The 5th picture also has a hat on the man's head obstructing part of his head, but it is still correctly classified.

A few improvements that could be made:

1. add image augmentation, so the algorithm will still perform well with shifted image.
2. potentially I could add more training set to increase accuracy.
3. potentially I could add more cnn layers to increase accuracy.

In [43]:


```
1  ## TODO: Execute your algorithm from Step 6 on
2  ## at least 6 images on your computer.
3  ## Feel free to use as many code cells as needed.
4
5  ## Load the cell
6  sample_files = np.array(glob("sample_pictures/*"))
7  print(sample_files)
8
9
```

```
['sample_pictures/images-3.jpeg' 'sample_pictures/images.jpeg'
'sample_pictures/images-2.jpeg'
'sample_pictures/images-2 \xe6\x8b\xb7\xe8\xb2\x9d.jpeg'
'sample_pictures/images \xe6\x8b\xb7\xe8\xb2\x9d 2.jpeg'
'sample_pictures/images \xe6\x8b\xb7\xe8\xb2\x9d.jpeg']
```

In [46]:




```
1 for path in sample_files:
2     dog_breed_detector(path)
```



That's a human in the image, but it looks like dog breed Maltese



That's a dog in the image. Breed: Bullmastiff



That's a human in the image, but it looks like dog breed Dachshund







That's a dog in the image. Breed: German\_shepherd\_dog



That's a human in the image, but it looks like dog breed Doberman\_pinscher



That's a dog in the image. Breed: American\_foxhound

In [ ]:

# Artificial Intelligence Nanodegree

## Convolutional Neural Networks

### Project: Write an Algorithm for a Dog Identification App

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all of the code implementations, you need to finalize your work by exporting the iPython Notebook as an HTML document. Before exporting the notebook to html, all of the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to "\n", "**File -> Download as -> HTML (.html)**". Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this IPython notebook.

---

## Why We're Here

In this notebook, you will make the first steps towards developing an algorithm that could be used as part of a mobile or web app. At the end of this project, your code will accept any user-supplied image as input. If a dog is detected in the image, it will provide an estimate of the dog's breed. If a human is detected, it will provide an estimate of the dog breed that is most resembling. The image below displays potential sample output of your finished project (... but we expect that each student's algorithm will behave differently!).



In this real-world setting, you will need to piece together a series of models to perform different tasks; for instance, the algorithm that detects humans in an image will be different from the CNN that infers dog breed. There are many points of possible failure, and no perfect algorithm exists. Your imperfect solution will nonetheless create a fun user experience!

## The Road Ahead

We break the notebook into separate steps. Feel free to use the links below to navigate the notebook.

- [Step 0](#): Import Datasets
- [Step 1](#): Detect Humans
- [Step 2](#): Detect Dogs
- [Step 3](#): Create a CNN to Classify Dog Breeds (from Scratch)
- [Step 4](#): Use a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 5](#): Create a CNN to Classify Dog Breeds (using Transfer Learning)
- [Step 6](#): Write your Algorithm
- [Step 7](#): Test Your Algorithm

---

## Step 0: Import Datasets

### Import Dog Dataset

In the code cell below, we import a dataset of dog images. We populate a few variables through the use of the `load_files` function from the scikit-learn library:

- `train_files`, `valid_files`, `test_files` - numpy arrays containing file paths to images
- `train_targets`, `valid_targets`, `test_targets` - numpy arrays containing onehot-encoded classification labels
- `dog_names` - list of string-valued dog breed names for translating labels

In [1]:

```
/Users/Anna/anaconda/envs/py27/lib/python2.7/site-packages/h5py/__init__
__.py:36: FutureWarning: Conversion of the second argument of issubdtype
from `float` to `np.floating` is deprecated. In future, it will be
treated as `np.float64 == np.dtype(float).type`.
```

```
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
There are 133 total dog categories.
There are 8351 total dog images.
```

```
There are 6680 training dog images.
There are 835 validation dog images.
There are 836 test dog images.
```

## Import Human Dataset

In the code cell below, we import a dataset of human images, where the file paths are stored in the numpy array `human_files`.

In [2]:

```
There are 13233 total human images.
```


# Step 1: Detect Humans

We use OpenCV's implementation of [Haar feature-based cascade classifiers](http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html) ([http://docs.opencv.org/trunk/d7/d8b/tutorial\\_py\\_face\\_detection.html](http://docs.opencv.org/trunk/d7/d8b/tutorial_py_face_detection.html)) to detect human faces in images. OpenCV provides many pre-trained face detectors, stored as XML files on [github](https://github.com/opencv/opencv/tree/master/data/haarcascades) (<https://github.com/opencv/opencv/tree/master/data/haarcascades>). We have downloaded one of these detectors and stored it in the `haarcascades` directory.

In the next code cell, we demonstrate how to use this detector to find human faces in a sample image.

In [3]:

```
('Number of faces detected:', 1)
```



Before using any of the face detectors, it is standard procedure to convert the images to grayscale. The `detectMultiScale` function executes the classifier stored in `face_cascade` and takes the grayscale image as a parameter.

In the above code, `faces` is a numpy array of detected faces, where each row corresponds to a detected face. Each detected face is a 1D array with four entries that specifies the bounding box of the detected face. The first two entries in the array (extracted in the above code as `x` and `y`) specify the horizontal and vertical positions of the top left corner of the bounding box. The last two entries in the array (extracted here as `w` and `h`) specify the width and height of the box.

## Write a Human Face Detector

We can use this procedure to write a function that returns `True` if a human face is detected in an image and `False` otherwise. This function, aptly named `face_detector`, takes a string-valued file path to an image as input and appears in the code block below.

In [4]:

## (IMPLEMENTATION) Assess the Human Face Detector

**Question 1:** Use the code cell below to test the performance of the `face_detector` function.

- What percentage of the first 100 images in `human_files` have a detected human face?
- What percentage of the first 100 images in `dog_files` have a detected human face?

Ideally, we would like 100% of human images with a detected face and 0% of dog images with a detected face. You will see that our algorithm falls short of this goal, but still gives acceptable performance. We extract the file paths for the first 100 images from each of the datasets and store them in the numpy arrays `human_files_short` and `dog_files_short`.

**Answer:**

Humans face detection accuracy: 99%

Dogs face detection accuracy: 11%

In [5]:

```
('Humans face detection accuracy:', 0.99)
('Dogs face detection accuracy:', 0.11)
```

**Question 2:** This algorithmic choice necessitates that we communicate to the user that we accept human images only when they provide a clear view of a face (otherwise, we risk having unnecessarily frustrated users!). In your opinion, is this a reasonable expectation to pose on the user? If not, can you think of a way to detect humans in images that does not necessitate an image with a clearly presented face?

**Answer:**

Yes, I think it is a reasonable expectation to pose to the user that we accept human images only when they provide a clear view of a face. In image detection task, there are many features OpenCV could be looking for to classify it as human. For example the boundaries of a human face, it would be hard to correctly classify it if the image of the human face is blurry or covered. The higher the quality of the image helps the algorithm to be more efficient with higher accuracy.

We suggest the face detector from OpenCV as a potential way to detect human images in your algorithm, but you are free to explore other approaches, especially approaches that make use of deep learning :). Please use the code cell below to design and test your own face detection algorithm. If you decide to pursue this *optional* task, report performance on each of the datasets.



In [6]:

## Step 2: Detect Dogs

In this section, we use a pre-trained [ResNet-50](#) (<http://ethereon.github.io/netscope/#/gist/db945b393d40bfa26006>) model to detect dogs in images. Our first line of code downloads the ResNet-50 model, along with weights that have been trained on [ImageNet](#) (<http://www.image-net.org/>), a very large, very popular dataset used for image classification and other vision tasks. ImageNet contains over 10 million URLs, each linking to an image containing an object from one of [1000 categories](#) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>). Given an image, this pre-trained ResNet-50 model returns a prediction (derived from the available categories in ImageNet) for the object that is contained in the image.

In [7]:

## Pre-process the Data

When using TensorFlow as backend, Keras CNNs require a 4D array (which we'll also refer to as a 4D tensor) as input, with shape

$$(nb\_samples, rows, columns, channels),$$

where `nb_samples` corresponds to the total number of images (or samples), and `rows`, `columns`, and `channels` correspond to the number of rows, columns, and channels for each image, respectively.

The `path_to_tensor` function below takes a string-valued file path to a color image as input and returns a 4D tensor suitable for supplying to a Keras CNN. The function first loads the image and resizes it to a square image that is  $224 \times 224$  pixels. Next, the image is converted to an array, which is then resized to a 4D tensor. In this case, since we are working with color images, each image has three channels. Likewise, since we are processing a single image (or sample), the returned tensor will always have shape

$$(1, 224, 224, 3).$$

The `paths_to_tensor` function takes a numpy array of string-valued image paths as input and returns a 4D tensor with shape

$$(nb\_samples, 224, 224, 3).$$

Here, `nb_samples` is the number of samples, or number of images, in the supplied array of image paths. It is best to think of `nb_samples` as the number of 3D tensors (where each 3D tensor corresponds to a different image) in your dataset!



In [8]:

## Making Predictions with ResNet-50

Getting the 4D tensor ready for ResNet-50, and for any other pre-trained model in Keras, requires some additional processing. First, the RGB image is converted to BGR by reordering the channels. All pre-trained models have the additional normalization step that the mean pixel (expressed in RGB as [103.939, 116.779, 123.68] and calculated from all pixels in all images in ImageNet) must be subtracted from every pixel in each image. This is implemented in the imported function `preprocess_input`. If you're curious, you can check the code for `preprocess_input` [here](https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py) ([https://github.com/fchollet/keras/blob/master/keras/applications/imagenet\\_utils.py](https://github.com/fchollet/keras/blob/master/keras/applications/imagenet_utils.py)).

Now that we have a way to format our image for supplying to ResNet-50, we are now ready to use the model to extract the predictions. This is accomplished with the `predict` method, which returns an array whose  $i$ -th entry is the model's predicted probability that the image belongs to the  $i$ -th ImageNet category. This is implemented in the `ResNet50_predict_labels` function below.

By taking the argmax of the predicted probability vector, we obtain an integer corresponding to the model's predicted object class, which we can identify with an object category through the use of this [dictionary](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>).

In [9]:

## Write a Dog Detector

While looking at the [dictionary](https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a) (<https://gist.github.com/yrevar/942d3a0ac09ec9e5eb3a>), you will notice that the categories corresponding to dogs appear in an uninterrupted sequence and correspond to dictionary keys 151-268, inclusive, to include all categories from 'Chihuahua' to 'Mexican hairless'. Thus, in order to check to see if an image is predicted to contain a dog by the pre-trained ResNet-50 model, we need only check if the `ResNet50_predict_labels` function above returns a value between 151 and 268 (inclusive).

We use these ideas to complete the `dog_detector` function below, which returns `True` if a dog is detected in an image (and `False` if not).

In [10]:

# (IMPLEMENTATION) Assess the Dog Detector

**Question 3:** Use the code cell below to test the performance of your `dog_detector` function.

- What percentage of the images in `human_files_short` have a detected dog?
- What percentage of the images in `dog_files_short` have a detected dog?

**Answer:**

Humans face detection accuracy: 2%

Dogs face detection accuracy: 100%

In [11]:

```
('Humans face detection accuracy:', 0.02)
('Dogs face detection accuracy:', 1.0)
```

## Step 3: Create a CNN to Classify Dog Breeds (from Scratch)

Now that we have functions for detecting humans and dogs in images, we need a way to predict breed from images. In this step, you will create a CNN that classifies dog breeds. You must create your CNN *from scratch* (so, you can't use transfer learning yet!), and you must attain a test accuracy of at least 1%. In Step 5 of this notebook, you will have the opportunity to use transfer learning to create a CNN that attains greatly improved accuracy.

Be careful with adding too many trainable layers! More parameters means longer training, which means you are more likely to need a GPU to accelerate the training process. Thankfully, Keras provides a handy estimate of the time that each epoch is likely to take; you can extrapolate this estimate to figure out how long it will take for your algorithm to train.

We mention that the task of assigning breed to dogs from images is considered exceptionally challenging. To see why, consider that *even a human* would have great difficulty in distinguishing between a Brittany and a Welsh Springer Spaniel.


| Brittany | Welsh Springer Spaniel |
|----------|------------------------|
|          |                        |



It is not difficult to find other dog breed pairs with minimal inter-class variation (for instance, Curly-Coated Retrievers and American Water Spaniels).

**Curly-Coated Retriever**

**American Water Spaniel**




Likewise, recall that labradors come in yellow, chocolate, and black. Your vision-based algorithm will have to conquer this high intra-class variation to determine how to classify all of these different shades as the same breed.

**Yellow Labrador**

**Chocolate Labrador**

**Black Labrador**



We also mention that random chance presents an exceptionally low bar: setting aside the fact that the classes are slightly imbalanced, a random guess will provide a correct answer roughly 1 in 133 times, which corresponds to an accuracy of less than 1%.

Remember that the practice is far ahead of the theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

## Pre-process the Data

We rescale the images by dividing every pixel in every image by 255.

In [12]:

```
100% |██████████| 6680/6680 [02:33<00:00, 43.39it/s]
100% |██████████| 835/835 [00:17<00:00, 48.52it/s]
100% |██████████| 836/836 [00:17<00:00, 48.48it/s]
```

## (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
model.summary()
```

We have imported some Python modules to get you started, but feel free to import as many modules as you need. If you end up getting stuck, here's a hint that specifies a model that trains relatively fast on CPU and attains >1% test accuracy in 5 epochs:



| Layer (type)  | Output Shape         | Param # |
|---|----------------------|---------|
| conv2d_1 (Conv2D)                                   | (None, 223, 223, 16) | 208     |
| max_pooling2d_1 (MaxPooling2D)                      | (None, 111, 111, 16) | 0       |
| conv2d_2 (Conv2D)                                   | (None, 110, 110, 32) | 2080    |
| max_pooling2d_2 (MaxPooling2D)                      | (None, 55, 55, 32)   | 0       |
| conv2d_3 (Conv2D)                                   | (None, 54, 54, 64)   | 8256    |
| max_pooling2d_3 (MaxPooling2D)                      | (None, 27, 27, 64)   | 0       |
| global_average_pooling2d_1 (GlobalAveragePooling2D) | (None, 64)           | 0       |
| dense_1 (Dense)                                     | (None, 133)          | 8645    |
| Total params: 19,189.0                              |                      |         |
| Trainable params: 19,189.0                          |                      |         |
| Non-trainable params: 0.0                           |                      |         |



**Question 4:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. If you chose to use the hinted architecture above, describe why you think that CNN architecture should work well for the image classification task.

**Answer:**

I chose a Convolutional layer of 16 filters to handle the initial input of the Neural Network. I also chose the activation function relu. I followed this layer with a max pooling layer of 2 which is standard.

Added another convolutional layer that is double the size as the first layer to get more detailed features as passed through from the prior layer. Nothing was changed here as it was designed purposely to be the exact same.

Following the first 2 layers I changed the third convolutional layer to include 64 filters for more details. This way I could give the network more information to classify the dog breed successfully. With the rest remain the same.

To avoid overfitting, I added dropout layer followed by a flatten layer so I could then add some fully connected layers to the network and dropout layer again. Then rape up with a output layer with softmax function.

In [13]:

| Layer (type)                   | Output Shape         | Param # |
|--------------------------------|----------------------|---------|
| =====                          |                      |         |
| conv2d_1 (Conv2D)              | (None, 224, 224, 16) | 208     |
| max_pooling2d_2 (MaxPooling2D) | (None, 112, 112, 16) | 0       |
| conv2d_2 (Conv2D)              | (None, 112, 112, 32) | 2080    |
| max_pooling2d_3 (MaxPooling2D) | (None, 56, 56, 32)   | 0       |
| conv2d_3 (Conv2D)              | (None, 56, 56, 64)   | 8256    |
| max_pooling2d_4 (MaxPooling2D) | (None, 28, 28, 64)   | 0       |
| dropout_1 (Dropout)            | (None, 28, 28, 64)   | 0       |
| flatten_2 (Flatten)            | (None, 50176)        | 0       |
| Total params: 105000           |                      |         |

## Compile the Model

In [14]:

## (IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.



In [15]:

```
Train on 6680 samples, validate on 835 samples
Epoch 1/5
6660/6680 [=====>.] - ETA: 1s - loss: 5.0752 -
acc: 0.0207
Epoch 00001: val_loss improved from inf to 4.58885, saving model to sa
ved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 428s 64ms/step - loss: 5.
0724 - acc: 0.0208 - val_loss: 4.5888 - val_acc: 0.0359
Epoch 2/5
6660/6680 [=====>.] - ETA: 1s - loss: 4.3535 -
acc: 0.0590
Epoch 00002: val_loss improved from 4.58885 to 4.24946, saving model t
o saved_models/weights.best.from_scratch.hdf5
6680/6680 [=====] - 466s 70ms/step - loss: 4.
3521 - acc: 0.0594 - val_loss: 4.2495 - val_acc: 0.0647
Epoch 3/5
6660/6680 [=====>.] - ETA: 1s - loss: 3.7908 -
acc: 0.1380
Epoch 00003: val_loss improved from 4.24946 to 4.09295, saving model t
```

## Load the Model with the Best Validation Loss

In [16]:

## Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 1%.

In [17]:

```
Test accuracy: 8.0000%
```

## Step 4: Use a CNN to Classify Dog Breeds

To reduce training time without sacrificing accuracy, we show you how to train a CNN using transfer learning. In the following step, you will get a chance to use transfer learning to train your own CNN.

## Obtain Bottleneck Features

In [18]:

## Model Architecture

The model uses the the pre-trained VGG-16 model as a fixed feature extractor, where the last convolutional output of VGG-16 is fed as input to our model. We only add a global average pooling layer and a fully connected layer, where the latter contains one node for each dog category and is equipped with a softmax.

In [19]:

| Layer (type)                 | Output Shape | Param # |
|------------------------------|--------------|---------|
| =====                        |              |         |
| global_average_pooling2d_1 ( | (None, 512)  | 0       |
| dense_3 (Dense)              | (None, 133)  | 68229   |
| =====                        |              |         |
| Total params: 68,229         |              |         |
| Trainable params: 68,229     |              |         |
| Non-trainable params: 0      |              |         |

## Compile the Model

In [20]:

## Train the Model

In [21]:

```
Train on 6680 samples, validate on 835 samples
Epoch 1/20
6580/6680 [=====>.] - ETA: 0s - loss: 12.3766 -
acc: 0.1309
Epoch 00001: val_loss improved from inf to 10.68090, saving model to s
aved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 4s 633us/step - loss: 12.
3627 - acc: 0.1316 - val_loss: 10.6809 - val_acc: 0.2287
Epoch 2/20
6540/6680 [=====>.] - ETA: 0s - loss: 10.1996 -
acc: 0.2872
Epoch 00002: val_loss improved from 10.68090 to 9.86518, saving model
to saved_models/weights.best.VGG16.hdf5
6680/6680 [=====] - 2s 261us/step - loss: 10.
1736 - acc: 0.2888 - val_loss: 9.8652 - val_acc: 0.3150
Epoch 3/20
6580/6680 [=====>.] - ETA: 0s - loss: 9.5918 -
acc: 0.3457
Epoch 00003: val_loss improved from 9.86518 to 9.58232, saving model t
o saved_models/weights.best.VGG16.hdf5
```

## Load the Model with the Best Validation Loss

In [22]:

## Test the Model

Now, we can use the CNN to test how well it identifies breed within our test dataset of dog images. We print the test accuracy below.

In [23]:

```
Test accuracy: 47.0000%
```

## Predict Dog Breed with the Model

In [24]:

## Step 5: Create a CNN to Classify Dog Breeds (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify dog breed from images. Your CNN must attain at least 60% accuracy on the test set.

In Step 4, we used transfer learning to create a CNN using VGG-16 bottleneck features. In this section, you must use the bottleneck features from a different pre-trained model. To make things easier for you, we have pre-computed the features for all of the networks that are currently available in Keras:

- [VGG-19 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogVGG19Data.npz) bottleneck features
- [ResNet-50 \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogResnet50Data.npz) bottleneck features
- [Inception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogInceptionV3Data.npz) bottleneck features
- [Xception \(https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz\)](https://s3-us-west-1.amazonaws.com/udacity-aind/dog-project/DogXceptionData.npz) bottleneck features

The files are encoded as such:

```
Dog{network}Data.npz
```

where `{network}`, in the above filename, can be one of `VGG19`, `Resnet50`, `InceptionV3`, or `Xception`. Pick one of the above architectures, download the corresponding bottleneck features, and store the downloaded file in the `bottleneck_features/` folder in the repository.

### (IMPLEMENTATION) Obtain Bottleneck Features

In the code block below, extract the bottleneck features corresponding to the train, test, and validation sets by running the following:

```
bottleneck_features = np.load('bottleneck_features/Dog{network}Data.npz')
train_{network} = bottleneck_features['train']
valid_{network} = bottleneck_features['valid']
test_{network} = bottleneck_features['test']
```

In [25]:

## (IMPLEMENTATION) Model Architecture

Create a CNN to classify dog breed. At the end of your code cell block, summarize the layers of your model by executing the line:

```
<your model's name>.summary()
```

**Question 5:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:**

I decided to leverage the current resnet weights and make that an input into the global average pooling layer to take advantage of transfer learning. So I could make use of what the network already understood from previous training. I added the fully connected layer with an output of 133 units, because there are 133 dognames, with the softmax function.

In [26]:

| Layer (type)               | Output Shape | Param # |
|----------------------------|--------------|---------|
| global_average_pooling2d_2 | (None, 2048) | 0       |
| dense_4 (Dense)            | (None, 133)  | 272517  |
| Total params: 272,517      |              |         |
| Trainable params: 272,517  |              |         |
| Non-trainable params: 0    |              |         |

## (IMPLEMENTATION) Compile the Model

In [27]:

## (IMPLEMENTATION) Train the Model

Train your model in the code cell below. Use model checkpointing to save the model that attains the best validation loss.

You are welcome to [augment the training data \(https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html\)](https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html), but this is not a requirement.

In [28]:

```
Train on 6680 samples, validate on 835 samples
Epoch 1/10
6560/6680 [=====>.] - ETA: 0s - loss: 1.6179 -
acc: 0.6003
Epoch 00001: val_loss improved from inf to 0.83293, saving model to sa
ved_models/weights.best.ResNet50.hdf5
6680/6680 [=====] - 3s 499us/step - loss: 1.6
011 - acc: 0.6037 - val_loss: 0.8329 - val_acc: 0.7425
Epoch 2/10
6660/6680 [=====>.] - ETA: 0s - loss: 0.4435 -
acc: 0.8614
Epoch 00002: val_loss improved from 0.83293 to 0.71584, saving model t
o saved_models/weights.best.ResNet50.hdf5
6680/6680 [=====] - 2s 348us/step - loss: 0.4
431 - acc: 0.8614 - val_loss: 0.7158 - val_acc: 0.7904
Epoch 3/10
6660/6680 [=====>.] - ETA: 0s - loss: 0.2663 -
acc: 0.9179
Epoch 00003: val_loss improved from 0.71584 to 0.66617, saving model t
```

## (IMPLEMENTATION) Load the Model with the Best Validation Loss

In [29]:

## (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of dog images. Ensure that your test accuracy is greater than 60%.

In [30]:

```
Test accuracy: 82.0000%
```



## (IMPLEMENTATION) Predict Dog Breed with the Model

Write a function that takes an image path as input and returns the dog breed ( `Affenpinscher` , `Afghan_hound` , etc) that is predicted by your model.

Similar to the analogous function in Step 5, your function should have three steps:

1. Extract the bottleneck features corresponding to the chosen CNN model.
2. Supply the bottleneck features as input to the model to return the predicted vector. Note that the `argmax` of this prediction vector gives the index of the predicted dog breed.
3. Use the `dog_names` array defined in Step 0 of this notebook to return the corresponding breed.

The functions to extract the bottleneck features can be found in `extract_bottleneck_features.py` , and they have been imported in an earlier code cell. To obtain the bottleneck features corresponding to your chosen CNN architecture, you need to use the function

```
extract_{network}
```

where `{network}` , in the above filename, should be one of `VGG19` , `Resnet50` , `InceptionV3` , or `Xception` .

In [31]:

## Step 6: Write your Algorithm


Write an algorithm that accepts a file path to an image and first determines whether the image contains a human, dog, or neither. Then,

- if a **dog** is detected in the image, return the predicted breed.
- if a **human** is detected in the image, return the resembling dog breed.
- if **neither** is detected in the image, provide output that indicates an error.

You are welcome to write your own functions for detecting humans and dogs in images, but feel free to use the `face_detector` and `dog_detector` functions developed above. You are **required** to use your CNN from Step 5 to predict dog breed.

Some sample output for our algorithm is provided below, but feel free to design your own user experience!

```
hello, human!
```



```
You look like a ...  
Chinese_shar-pei
```

## (IMPLEMENTATION) Write your Algorithm

In [32]:

```
A local file was found, but it seems to be incomplete or outdated because the md5 file hash does not match the original value of a268eb855778b3df3c7506639542a6af so we will re-download the data.
```

```
Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://github.com/fchollet/deep-learning-models/releases/download/v0.2/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5)
```

```
94658560/94653016 [=====] - 53s 1us/step
```

```
94666752/94653016 [=====] - 53s 1us/step
```



## Step 7: Test Your Algorithm

In this section, you will take your new algorithm for a spin! What kind of dog does the algorithm think that **you** look like? If you have a dog, does it predict your dog's breed accurately? If you have a cat, does it mistakenly think that your cat is a dog?

### (IMPLEMENTATION) Test Your Algorithm on Sample Images!

Test your algorithm at least six images on your computer. Feel free to use any images you like. Use at least two human and two dog images.

**Question 6:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

## Answer:

The output is better than I expected since all 6 images are correctly classified between human and dog, also dog breed are all correctly classified! Even in picture2 Bullmastiff were wearing clothes, the detector still correctly classified it. The 5th picture also has a hat on the man's head obstructing part of his head, but it is still correctly classified.


A few improvements that could be made:

1. add image augmentation, so the algorithm will still perform well with shifted image.
2. potentially I could add more training set to increase accuracy.
3. potentially I could add more cnn layers to increase accuracy.

In [43]:

```
['sample_pictures/images-3.jpeg' 'sample_pictures/images.jpeg'
'sample_pictures/images-2.jpeg'
'sample_pictures/images-2 \xe6\x8b\xb7\xe8\xb2\x9d.jpeg'
'sample_pictures/images \xe6\x8b\xb7\xe8\xb2\x9d 2.jpeg'
'sample_pictures/images \xe6\x8b\xb7\xe8\xb2\x9d.jpeg']
```

In [46]:



That's a human in the image, but it looks like dog breed Maltese



In [ ]: